

## CNN with tensorflow

### Convolutional Neural Network (CNN)

This tutorial demonstrates training a simple Convolutional Neural Network (CNN) to classify CIFAR images. Because this tutorial uses the Keras Sequential API, creating and training our model will take just a few lines of code.

Import TensorFlow

```
1 import tensorflow as tf
2 from tensorflow.keras import datasets, layers, models
3 import matplotlib.pyplot as plt
```

Download and prepare the CIFAR10 dataset

The CIFAR10 dataset contains 60,000 color images in 10 classes, with 6,000 images in each class. The dataset is divided into 50,000 training images and 10,000 testing images. The classes are mutually exclusive and there is no overlap between them.

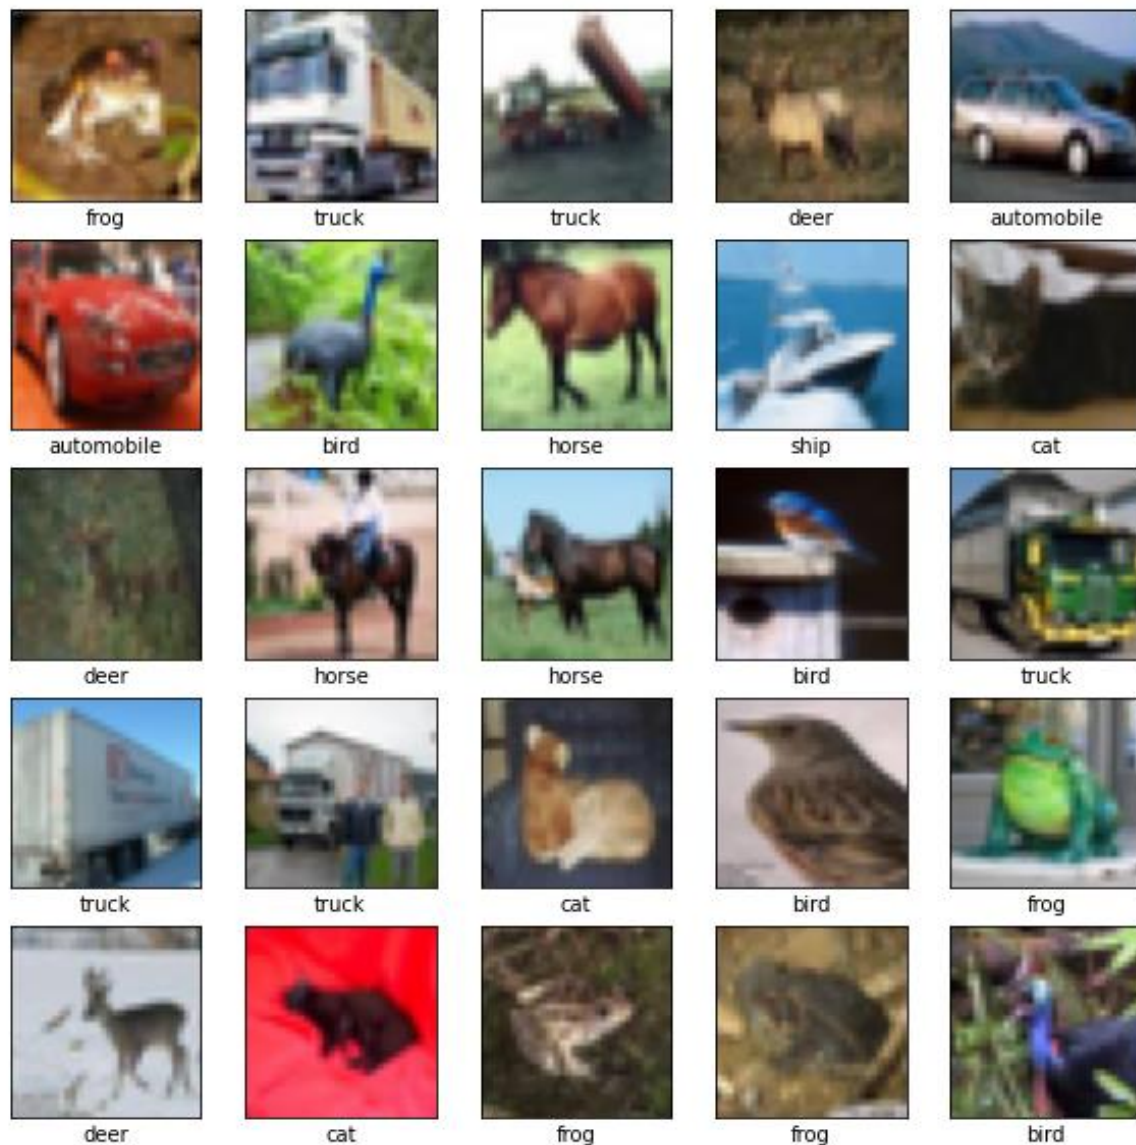
```
1 (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
2
3 # Normalize pixel values to be between 0 and 1
4 train_images, test_images = train_images / 255.0, test_images / 255.0
5
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170500096/170498071 [=====] - 6s 0us/step

Verify the data

To verify that the dataset looks correct, let's plot the first 25 images from the training set and display the class name below each image.

```
1 class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
2               'dog', 'frog', 'horse', 'ship', 'truck']
3
4 plt.figure(figsize=(10,10))
5 for i in range(25):
6     plt.subplot(5,5,i+1)
7     plt.xticks([])
8     plt.yticks([])
9     plt.grid(False)
10    plt.imshow(train_images[i], cmap=plt.cm.binary)
11    # The CIFAR labels happen to be arrays,
12    # which is why you need the extra index
13    plt.xlabel(class_names[train_labels[i][0]])
14 plt.show()
15
```



## Create the convolutional base

The 6 lines of code below define the convolutional base using a common pattern: a stack of Conv2D and MaxPooling2D layers.

As input, a CNN takes tensors of shape (image\_height, image\_width, color\_channels), ignoring the batch size. If you are new to these dimensions, color\_channels refers to (R,G,B). In this example, you will configure our CNN to process inputs of shape (32, 32, 3), which is the format of CIFAR images. You can do this by passing the argument `input_shape` to our first layer.

```
1 model = models.Sequential()
2 model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
3 model.add(layers.MaxPooling2D((2, 2)))
4 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
5 model.add(layers.MaxPooling2D((2, 2)))
6 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
7
```

Let's display the architecture of our model so far.

```
1 model.summary()  
2
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
Total params: 56,320		
Trainable params: 56,320		
Non-trainable params: 0		

Above, you can see that the output of every Conv2D and MaxPooling2D layer is a 3D tensor of shape (height, width, channels). The width and height dimensions tend to shrink as you go deeper in the network. The number of output channels for each Conv2D layer is controlled by the first argument (e.g., 32 or 64). Typically, as the width and height shrink, you can afford (computationally) to add more output channels in each Conv2D layer.

## Add Dense layers on top

To complete our model, you will feed the last output tensor from the convolutional base (of shape (4, 4, 64)) into one or more Dense layers to perform classification. Dense layers take vectors as input (which are 1D), while the current output is a 3D tensor. First, you will flatten (or unroll) the 3D output to 1D, then add one or more Dense layers on top. CIFAR has 10 output classes, so you use a final Dense layer with 10 outputs and a softmax activation.

```
1 model.add(layers.Flatten())  
2 model.add(layers.Dense(64, activation='relu'))  
3 model.add(layers.Dense(10))  
4
```

Here's the complete architecture of our model.

```
1 model.summary()
2
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 64)	65600
dense_1 (Dense)	(None, 10)	650
Total params: 122,570		
Trainable params: 122,570		
Non-trainable params: 0		

As you can see, our (4, 4, 64) outputs were flattened into vectors of shape (1024) before going through two Dense layers.

## Compile and train the model


```
1 model.compile(optimizer='adam',
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3               metrics=['accuracy'])
4
5 history = model.fit(train_images, train_labels, epochs=10,
6                     validation_data=(test_images, test_labels))
7
```

```
Epoch 1/10
1563/1563 [=====] - 65s 41ms/step - loss: 1.7444 - accuracy: 0.3514 - val_loss: 1.2495 - val_accuracy: 0.5474
Epoch 2/10
1563/1563 [=====] - 63s 41ms/step - loss: 1.2088 - accuracy: 0.5688 - val_loss: 1.0813 - val_accuracy: 0.6176
Epoch 3/10
1563/1563 [=====] - 62s 40ms/step - loss: 1.0274 - accuracy: 0.6373 - val_loss: 1.0251 - val_accuracy: 0.6359
Epoch 4/10
1563/1563 [=====] - 63s 40ms/step - loss: 0.9259 - accuracy: 0.6756 - val_loss: 0.9984 - val_accuracy: 0.6523
Epoch 5/10
1563/1563 [=====] - 62s 40ms/step - loss: 0.8645 - accuracy: 0.6961 - val_loss: 0.8952 - val_accuracy: 0.6927
Epoch 6/10
1563/1563 [=====] - 62s 39ms/step - loss: 0.8022 - accuracy: 0.7212 - val_loss: 0.9012 - val_accuracy: 0.6914
Epoch 7/10
1563/1563 [=====] - 62s 40ms/step - loss: 0.7402 - accuracy: 0.7418 - val_loss: 0.9034 - val_accuracy: 0.6904
Epoch 8/10
1563/1563 [=====] - 62s 40ms/step - loss: 0.6987 - accuracy: 0.7559 - val_loss: 0.8622 - val_accuracy: 0.7132
Epoch 9/10
1563/1563 [=====] - 63s 40ms/step - loss: 0.6546 - accuracy: 0.7686 - val_loss: 0.8790 - val_accuracy: 0.7111
Epoch 10/10
1563/1563 [=====] - 63s 40ms/step - loss: 0.6188 - accuracy: 0.7832 - val_loss: 0.9414 - val_accuracy: 0.6945
```

## Evaluate the model

```
1 plt.plot(history.history['accuracy'], label='accuracy')
2 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
3 plt.xlabel('Epoch')
4 plt.ylabel('Accuracy')
5 plt.ylim([0.5, 1])
6 plt.legend(loc='lower right')
7
8 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
9
```

313/313 - 3s - loss: 0.9414 - accuracy: 0.6945



```
1 print(test_acc)
2
```

0.6945000290870667

## Image classification

This tutorial shows how to classify images of flowers. It creates an image classifier using a `keras.Sequential` model, and loads data using `preprocessing.image_dataset_from_directory`. You will gain practical experience with the following concepts:

- Efficiently loading a dataset off disk.
- Identifying overfitting and applying techniques to mitigate it, including data augmentation and Dropout.

This tutorial follows a basic machine learning workflow:

1. Examine and understand data
2. Build an input pipeline

3. Build the model
4. Train the model
5. Test the model
6. Improve the model and repeat the process

## Import TensorFlow and other libraries

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import os
4 import PIL
5 import tensorflow as tf
6
7 from tensorflow import keras
8 from tensorflow.keras import layers
9 from tensorflow.keras.models import Sequential
10
```

## Download and explore the dataset

This tutorial uses a dataset of about 3,700 photos of flowers. The dataset contains 5 sub-directories, one per class:

```
1 flower_photo/
2   daisy/
3   dandelion/
4   roses/
5   sunflowers/
6   tulips/
7
```

```
1 import pathlib
2 dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
3 data_dir = tf.keras.utils.get_file('flower_photos', origin=dataset_url, untar=True)
4 data_dir = pathlib.Path(data_dir)
5
```

Downloading data from [https://storage.googleapis.com/download.tensorflow.org/example\\_images/flower\\_photos.tgz](https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz)  
228818944/228813984 [=====] - 4s 0us/step

After downloading, you should now have a copy of the dataset available. There are 3,670 total images

```
1 image_count = len(list(data_dir.glob('*/*.jpg')))
2 print(image_count)
3
```

3670



Here are some roses:

```
1 roses = list(data_dir.glob('roses/*'))  
2 PIL.Image.open(str(roses[0]))  
3 |
```



```
1 PIL.Image.open(str(roses[1]))  
2
```



And some tulips:

```
1 tulips = list(data_dir.glob('tulips/*'))  
2 PIL.Image.open(str(tulips[0]))  
3
```



## Load using keras.preprocessing

Let's load these images off disk using the helpful `image_dataset_from_directory` utility. This will take you from a directory of images on disk to a `tf.data.Dataset` in just a couple lines of code. If you like, you can also write your own data loading code from scratch

### Create a dataset

Define some parameters for the loader:

```
1 batch_size = 32
2 img_height = 180
3 img_width = 180
4
```

It's good practice to use a validation split when developing your model. Let's use 80% of the images for training, and 20% for validation.

### Training set

```
1 train_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     data_dir,
3     validation_split=0.2,
4     subset="training",
5     seed=123,
6     image_size=(img_height, img_width),
7     batch_size=batch_size)
8
```

```
Found 3670 files belonging to 5 classes.
Using 2936 files for training.
```

### Validation set

```
1 val_ds = tf.keras.preprocessing.image_dataset_from_directory(
2     data_dir,
3     validation_split=0.2,
4     subset="validation",
5     seed=123,
6     image_size=(img_height, img_width),
7     batch_size=batch_size)
8
```

```
Found 3670 files belonging to 5 classes.
Using 734 files for validation.
```



You can find the class names in the `class_names` attribute on these datasets. These correspond to the directory names in alphabetical order.

```
1 class_names = train_ds.class_names
2 print(class_names)
3
4
```

```
['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
```

---

## Visualize the data

Here are the first 9 images from the training dataset.

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(10, 10))
4 for images, labels in train_ds.take(1):
5     for i in range(9):
6         ax = plt.subplot(3, 3, i + 1)
7         plt.imshow(images[i].numpy().astype("uint8"))
8         plt.title(class_names[labels[i]])
9         plt.axis("off")
10
```



You will train a model using these datasets by passing them to `model.fit` in a moment. If you like, you can also manually iterate over the dataset and retrieve batches of images:

The `image_batch` is a tensor of the shape `(32, 180, 180, 3)`. This is a batch of 32 images of shape `180x180x3` (the last dimension refers to color channels RGB). The `label_batch` is a tensor of the shape `(32,)`, these are corresponding labels to the 32 images.

You can call `.numpy()` on the `image_batch` and `labels_batch` tensors to convert them to a `numpy.ndarray`.

## Configure the dataset for performance

Let's make sure to use buffered prefetching so you can yield data from disk without having I/O become blocking. These are two important methods you should use when loading data.

`Dataset.cache()` keeps the images in memory after they're loaded off disk during the first epoch. This will ensure the dataset does not become a bottleneck while training your model. If your dataset is too large to fit into memory, you can also use this method to create a performant on-disk cache.

`Dataset.prefetch()` overlaps data preprocessing and model execution while training.

Interested readers can learn more about both methods, as well as how to cache data to disk in the data performance guide.

```
1 AUTOTUNE = tf.data.experimental.AUTOTUNE
2
3 train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
4 val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
5
```

## Standardize the data

The RGB channel values are in the `[0, 255]` range. This is not ideal for a neural network; in general you should seek to make your input values small. Here, you will standardize values to be in the `[0, 1]` range by using a Rescaling layer.

```
1 normalization_layer = layers.experimental.preprocessing.Rescaling(1./255)
2
```

There are two ways to use this layer. You can apply it to the dataset by calling `map`:

```
1 normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
2 image_batch, labels_batch = next(iter(normalized_ds))
3 first_image = image_batch[0]
4 # Notice the pixels values are now in `[0,1]`.
5 print(np.min(first_image), np.max(first_image))
6
```

```
0.0 0.9870374
```

Or, you can include the layer inside your model definition, which can simplify deployment. We'll see second approach later.

## Create the model

The model consists of three convolution blocks with a max pool layer in each of them. There's a fully connected layer with 128 units on top of it that is activated by a `relu` activation function. This model has not been tuned for high accuracy, the goal of this tutorial is to show a standard approach

```

1 num_classes = 5
2
3 model = Sequential([
4     layers.experimental.preprocessing.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
5     layers.Conv2D(16, 3, padding='same', activation='relu'),
6     layers.MaxPooling2D(),
7     layers.Conv2D(32, 3, padding='same', activation='relu'),
8     layers.MaxPooling2D(),
9     layers.Conv2D(64, 3, padding='same', activation='relu'),
10    layers.MaxPooling2D(),
11    layers.Flatten(),
12    layers.Dense(128, activation='relu'),
13    layers.Dense(num_classes)
14 ])
15

```

## Compile the model

For this tutorial, choose the `optimizers.Adam` optimizer and `losses.SparseCategoricalCrossentropy` loss function. To view training and validation accuracy for each training epoch, pass the `metrics` argument.

```

1 model.compile(optimizer='adam',
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3               metrics=['accuracy'])
4

```

## Model summary

View all the layers of the network using the model's `summary` method:

```
1 model.summary()
2
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_2 (MaxPooling2)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_3 (MaxPooling2)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_4 (MaxPooling2)	(None, 22, 22, 64)	0
flatten_1 (Flatten)	(None, 30976)	0
dense_2 (Dense)	(None, 128)	3965056
dense_3 (Dense)	(None, 5)	645
Total params: 3,989,285		
Trainable params: 3,989,285		
Non-trainable params: 0		

## Train the model

```
1 epochs=10
2 history = model.fit(
3     train_ds,
4     validation_data=val_ds,
5     epochs=epochs
6 )
7
```


```
Epoch 1/10
92/92 [=====] - 85s 918ms/step - loss: 1.6004 - accuracy: 0.3601 - val_loss: 1.0398 - val_accuracy: 0.5899
Epoch 2/10
92/92 [=====] - 82s 896ms/step - loss: 0.9737 - accuracy: 0.6153 - val_loss: 0.9661 - val_accuracy: 0.6158
Epoch 3/10
92/92 [=====] - 82s 896ms/step - loss: 0.8063 - accuracy: 0.6982 - val_loss: 0.9045 - val_accuracy: 0.6417
Epoch 4/10
92/92 [=====] - 83s 899ms/step - loss: 0.5645 - accuracy: 0.8061 - val_loss: 0.8798 - val_accuracy: 0.6676
Epoch 5/10
92/92 [=====] - 83s 902ms/step - loss: 0.3209 - accuracy: 0.8996 - val_loss: 1.0014 - val_accuracy: 0.6376
Epoch 6/10
92/92 [=====] - 83s 902ms/step - loss: 0.1637 - accuracy: 0.9577 - val_loss: 1.2816 - val_accuracy: 0.6185
Epoch 7/10
92/92 [=====] - 83s 903ms/step - loss: 0.1137 - accuracy: 0.9672 - val_loss: 1.2762 - val_accuracy: 0.6403
Epoch 8/10
92/92 [=====] - 83s 904ms/step - loss: 0.0525 - accuracy: 0.9913 - val_loss: 1.7757 - val_accuracy: 0.6471
Epoch 9/10
92/92 [=====] - 83s 904ms/step - loss: 0.0468 - accuracy: 0.9905 - val_loss: 1.6880 - val_accuracy: 0.6322
Epoch 10/10
92/92 [=====] - 83s 900ms/step - loss: 0.0302 - accuracy: 0.9934 - val_loss: 1.8118 - val_accuracy: 0.6417
```

## Visualize training results

Create plots of loss and accuracy on the training and validation sets.

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6
7 epochs_range = range(epochs)
8
9 plt.figure(figsize=(8, 8))
10 plt.subplot(1, 2, 1)
11 plt.plot(epochs_range, acc, label='Training Accuracy')
12 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
13 plt.legend(loc='lower right')
14 plt.title('Training and Validation Accuracy')
15
16 plt.subplot(1, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```





As you can see from the plots, training accuracy and validation accuracy are off by large margin and the model has achieved only around 60% accuracy on the validation set.

Let's look at what went wrong and try to increase the overall performance of the model.

## Data augmentation

Overfitting generally occurs when there are a small number of training examples. Data augmentation takes the approach of generating additional training data from your existing examples by augmenting them using random transformations that yield believable-looking images. This helps expose the model to more aspects of the data and generalize better.

You will implement data augmentation using experimental Keras Preprocessing Layers. These can be included inside your model like other layers, and run on the GPU.

Let's visualize what a few augmented examples look like by applying data augmentation to the same image several times:

You will use data augmentation to train a model in a moment.

## Dropout

Another technique to reduce overfitting is to introduce Dropout to the network, a form of *regularization*.

When you apply Dropout to a layer it randomly drops out (by setting the activation to zero) a number of output units from the layer during the training process. Dropout takes a fractional number as its input value, in the form such as 0.1, 0.2, 0.4, etc. This means dropping out 10%, 20% or 40% of the output units randomly from the applied layer.

Let's create a new neural network using `layers.Dropout`, then train it using augmented images.

```
1 model = Sequential([
2     data_augmentation,
3     layers.experimental.preprocessing.Rescaling(1./255),
4     layers.Conv2D(16, 3, padding='same', activation='relu'),
5     layers.MaxPooling2D(),
6     layers.Conv2D(32, 3, padding='same', activation='relu'),
7     layers.MaxPooling2D(),
8     layers.Conv2D(64, 3, padding='same', activation='relu'),
9     layers.MaxPooling2D(),
10    layers.Dropout(0.2),
11    layers.Flatten(),
12    layers.Dense(128, activation='relu'),
13    layers.Dense(num_classes)
14 ])
15
```

## Compile and train the model

```
1 model.compile(optimizer='adam',
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
3               metrics=['accuracy'])
4
```

```
1 model.summary()
2
```


Model: "sequential\_3"

Layer (type)	Output Shape	Param #
sequential_2 (Sequential)	(None, 180, 180, 3)	0
rescaling_2 (Rescaling)	(None, 180, 180, 3)	0
conv2d_6 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_5 (MaxPooling2)	(None, 90, 90, 16)	0
conv2d_7 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_6 (MaxPooling2)	(None, 45, 45, 32)	0
conv2d_8 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_7 (MaxPooling2)	(None, 22, 22, 64)	0
dropout (Dropout)	(None, 22, 22, 64)	0
flatten_2 (Flatten)	(None, 30976)	0
dense_4 (Dense)	(None, 128)	3965056
dense_5 (Dense)	(None, 5)	645

Total params: 3,989,285  
Trainable params: 3,989,285  
Non-trainable params: 0

## Visualize training results

After applying data augmentation and Dropout, there is less overfitting than before, and training and validation accuracy are closer aligned.



## Predict on new data

Finally, let's use our model to classify an image that wasn't included in the training or validation sets

```
1 sunflower_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg"
2 sunflower_path = tf.keras.utils.get_file('Red_sunflower', origin=sunflower_url)
3
4 img = keras.preprocessing.image.load_img(
5     sunflower_path, target_size=(img_height, img_width)
6 )
7 img_array = keras.preprocessing.image.img_to_array(img)
8 img_array = tf.expand_dims(img_array, 0) # Create a batch
9
10 predictions = model.predict(img_array)
11 score = tf.nn.softmax(predictions[0])
12
13 print(
14     "This image most likely belongs to {} with a {:.2f} percent confidence."
15     .format(class_names[np.argmax(score)], 100 * np.max(score))
16 )
17 |
```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/example\_images/592px-Red\_sunflower.jpg  
122880/117948 [=====] - 0s 0us/step  
This image most likely belongs to sunflowers with a 99.45 percent confidence.